

Deadlock (I)

The Deadlock Problem

❏ A set of blocked processes each holds an instance of resource and waits to acquire an instance of resource held by another process

❏ Example

❏ System has 2 disk drives

❏ P_1 and P_2 each hold one disk drive and each needs another one

❏ Example

❏ semaphores A and B , initialized to 1

P_0	P_1
wait (A);	wait(B)
wait (B);	wait(A)

Deadlock Characterization

Deadlock arises → 4 conditions hold simultaneously
(necessary conditions of deadlock)

- ❏ **Mutual exclusion:** an instance of resource can only be used by a process at a time
- ❏ **Hold and wait:** a process holding at least one instance of resource is waiting to acquire additional resources held by other processes
- ❏ **No preemption:** resource can be released only voluntarily by the process holding it, after that process has completed its task
- ❏ **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource in the following order:
request \rightarrow use \rightarrow release

Resource-Allocation Graph

A set of vertices V and a set of edges E .

☐ V is partitioned into two types:

☐ $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system


☐ $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system

☐ **request edge** – directed edge $P_i \rightarrow R_j$


☐ **assignment edge** – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


- Process
- Resource Type with 4 instances




- P_i requests an instance of R_j




- P_i is holding an instance of R_j



Example of a Resource Allocation Graph




Resource Allocation Graph With A Deadlock



There is one or more cycles in the graph

Graph With A Cycle But No Deadlock



Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock

Methods for Handling Deadlocks

- ❏ Ensure that the system will *never* enter a deadlock state
- ❏ Allow the system to enter a deadlock state and then recover
- ❏ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

Deadlock Prevention

Restrain the ways that request can be made such that at least one of the deadlock necessary conditions fails

- ❏ **Mutual Exclusion** → **Make each instance of resource simultaneously sharing** – often impossible
- ❏ **Hold and Wait** → **Request all at once; or, Give up all before request** – must guarantee that whenever a process requests a resource, it does not hold any other resources (i.e., it should release all resources that it currently holds)
 - ❏ Low resource utilization; starvation possible

Deadlock Prevention

- ❏ **No Preemption** → **Allow resource instance to be re-assignable**
 - ❏ Often impossible
- ❏ **Circular Wait** → **Prevent Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

Deadlock due to circular wait

P1 & P2 share semaphores A & B & C

P1:

wait(A)

...segment 1...

wait(B)

...segment 2...

wait(C)

...segment 3 ...

signal(C)

signal(B)

signal(A)

P2:

wait(B)

...segment 4...

wait(C)

...segment 5...

wait(A)

...segment 6 ...

signal(A)

signal(C)

signal(B)

Prevent circular wait

P1 & P2 share semaphores A & B & C; impose order: $A < B < C$

P1:

wait(A)

...segment 1...

wait(B)

...segment 2...

wait(C)

...segment 3 ...

signal(C)

signal(B)

signal(A)

P2:

wait(A)

wait(B)

...segment 4...

wait(C)

...segment 5...

...segment 6 ...

signal(A)

signal(C)

signal(B)