

Client Server Computing

- Although the Internet provides a basic communication service, the protocol software cannot initiate contact with, or accept contact from, a remote computer. Instead, two application programs must participate in any communication with one application initiates communication and the one accepts it.
- In network applications, a server application waits passively for contact after informing local protocol software that a specific type of message is expected, while a client application initiates communication actively by sending a matched type of message.

Characteristics of Clients Software

- is an arbitrary application program that becomes a client temporarily when remote access is needed, but also performs other computation locally.
- is invoked directly by a user, and executes only for one session.
- runs locally on a user's personal computer.
- actively initiates contact with a server.
- can access multiple services as needed, but actively contacts one remote server at a time.
- does not require special hardware or sophisticated operating system.

Characteristics of Servers Software

- is a special-purpose, privileged program dedicated to providing one service, but can handle multiple remote clients at the same time.
- is invoked automatically when a system boots, and continues to execute through many sessions.
- runs on a shared computer (i.e., not on a user's personal computer).
- waits passively for contact from arbitrary clients.
- accepts contact from arbitrary clients, but offers a single service.
- requires powerful hardware and a sophisticated OS.

Identifying A Particular Service

- Transport protocols assign each service a unique identifier.
- Both client and server specify the service identifier; protocol software uses the identifier to direct each incoming request to the correct server.
- In TCP/IP, TCP uses 16-bit integer values known as **protocol port numbers** to identify services.

Concurrent Server

- Concurrent execution is fundamental to servers because concurrency permits multiple clients to obtain a given service without having to wait for the server to finish previous requests.
- In concurrent server designs, the server creates a new thread to handle each client.
- Transport protocols assign an identifier to each client as well as to each service.
- Protocol software on the server's machine uses the combination of client and server identifiers to choose the correct copy of a concurrent server.


The Socket API

- The interface between an application program and the communication protocols in an operating system (OS) is known as the Application Program Interface or API.
- Sockets provide an implementation of the SAP (Service Access Point) abstraction at the Transport Layer in the TCP/IP protocol suite, which is part of the BSD Unix.
- A socket library can provide applications with a socket API on an operating system that does not provide native sockets (e.g. Windows 3.1). When an application calls one of the socket procedures, control passes to a library routine that makes one or more calls to the underlying OS to implement the socket function.

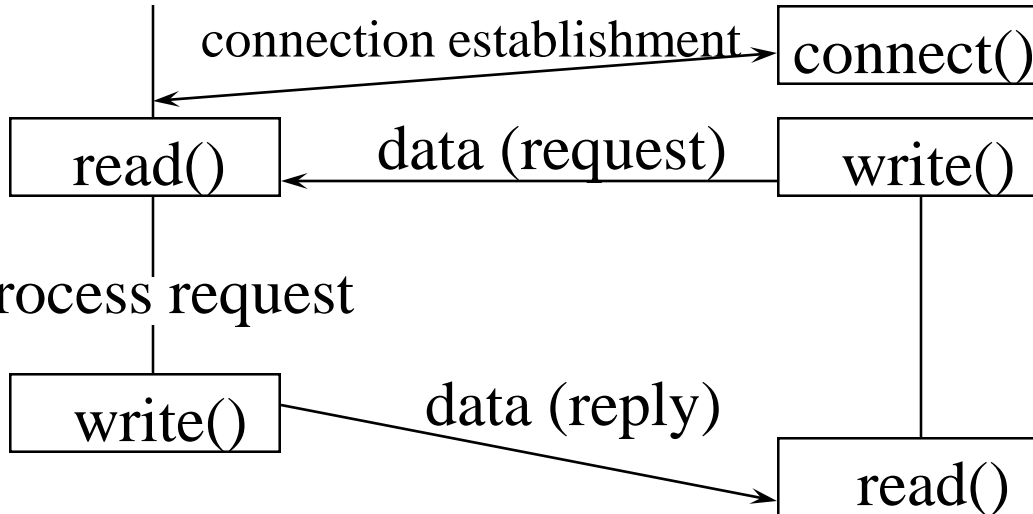
- A socket may be thought of as a generalization of the BSD Unix file access mechanism (open-read-write-close) that provides an end-point for communication.
- When an application creates a socket, the application is given a small integer descriptor used to reference the socket. If a system uses the same descriptor space for sockets and other I/O, a single application can be used for network communication as well as for local data transfer.
- An application must supply many details for each socket by specifying many parameters and options (e.g. an application must choose a particular protocol, provide address of remote machine, specify whether it is a client or server, etc.)
- To avoid having a single socket function with separate parameters for each options, designers of the socket API chose to define many functions, each with a few parameters.

Server

(connection-oriented protocol)




blocks until connection
from client



process request

Socket system calls for connection-oriented protocol

Client



Server

(connectionless protocol)

socket()

bind()

recvfrom()

blocks until data
received from client

process request

sendto()

Socket system calls for connectionless protocol

Client

socket()

bind()

sendto()

recvfrom()

data (request)

data (reply)

- Data communication between two hosts on the Internet require the five components of what is called an **association** to be initialized: {*protocol, local-addr, local-process, foreign-addr, foreign-process*}
- The different system calls for sockets provides values for one or more of these components.

Socket system call

- The first system call any process wishing to do network I/O has to call is the **socket** system call.
 - `int listenfd = socket (int family, int type, int protocol)`
 - Examples of Family include:
 - `PF_UNIX`
 - `PF_INET`
 - Examples of Type include
 - `SOCK_STREAM`
 - `SOCK_DGRAM`
 - `SOCK_RAW`
 - The protocol argument is typically zero, but may be specified to request an actual protocol like UDP, TCP, ICMP, etc.
 - Ideally, the three parameters should be orthogonal, but in reality, not all combinations are meaningful.

- The socket system call just fills in one element of the five-tuple we've looked at - the protocol. The remaining are filled in by the other calls as shown in the figure.

	<i>protocol</i>	<i>local_addr, local_process</i>	<i>foreign_addr, foreign_process</i>
Connection-Oriented Server	socket()	bind()	accept()
Connection-oriented Client	socket()	connect()	
Connectionless Server	socket()	bind()	recvfrom()
Connectionless Client	socket()	bind()	sendto()

Bind System Call

- The bind system call assigns an address to an unnamed socket. Example
 - int **bind**(int *listenfd*, struct sockaddr_in **myaddr*, int *addrlen*)
 - What is *bind* used for ?
 - *Servers* (both connection oriented and connectionless) NEED to register their well-known address to be able to accept connection requests.
 - A *client* can register a specific address for itself.
 - A *connectionless client* NEEDS to assure that it is bound to some unique address, so that the server has a valid return address to send its responses to.

- The *bind* system call provides the values for the *local_addr* and *local_process* elements in the *five_tuple* in an association.
- An address for the Internet domain sockets is a combination of a hostname and a port number, as shown below:

- struct sockaddr_in {
 - short sin_family ; /*typically AF_INET*/
 - u_short sin_port; /* 16 bit port number, *network byte ordered* */
 - struct in_addr sin_addr ; /* 32 bit netid/hostid, *network byte ordered* */
 - char sin_zero[8]; /* unused*/
 - }

Connect/Listen/Accept System Calls

- Connect
 - A **client** process *connects* a socket descriptor after a *socket* system call to establish a connection with the server.
 - int **connect**(int *listenfd*, struct sockaddr_in **servaddr*, int *addrlen*)
 - For a *connection-oriented client*, the **connect** (along with an **accept** at the server side) assigns all four addresses and process components of the association.

- ***Listen***

- The listen system call is used by a connection-oriented server to indicate it is willing to receive connections, e.g., **listen**(*listenfd*, *qlength*), where the system will enqueue up to *qlength* requests for connections.

- ***Accept***

- After the *server* executes a *listen*, it waits for connection requests from client(s) in the ***accept*** system call, e.g., *connfd* = **accept**(*listenfd*, *peer*, *addrlen*)
 - **accept** returns a new socket descriptor, which has all five components of the association specified - three (*protocol*, *local addr*, *local_process*) are inherited from the existing *listenfd* (which has its foreign address and process components unspecified, and hence can be re-used to accept another request. This scenario is typical for concurrent servers.)

Sending and Receiving Data

- Here's how you might read from a socket:
 - `num_read = read(listenfd, buff_ptr, num_bytes)`
- And here's how you read from an open file descriptor in Unix:
 - `num_read = read(fildes, buff_ptr, numbytes)`
- There are other ways (with different parameters) to send and receive data: `read`, `readv`, `recv`, `recvfrom`, `recvmsg` to receive data through a socket; and `write`, `writv`, `send`, `sendto`, `sendmsg` to send data through a socket.




Figure 1. TCP concurrent server with a passive open on port 21.




Figure 2. Connection request from client to concurrent server.

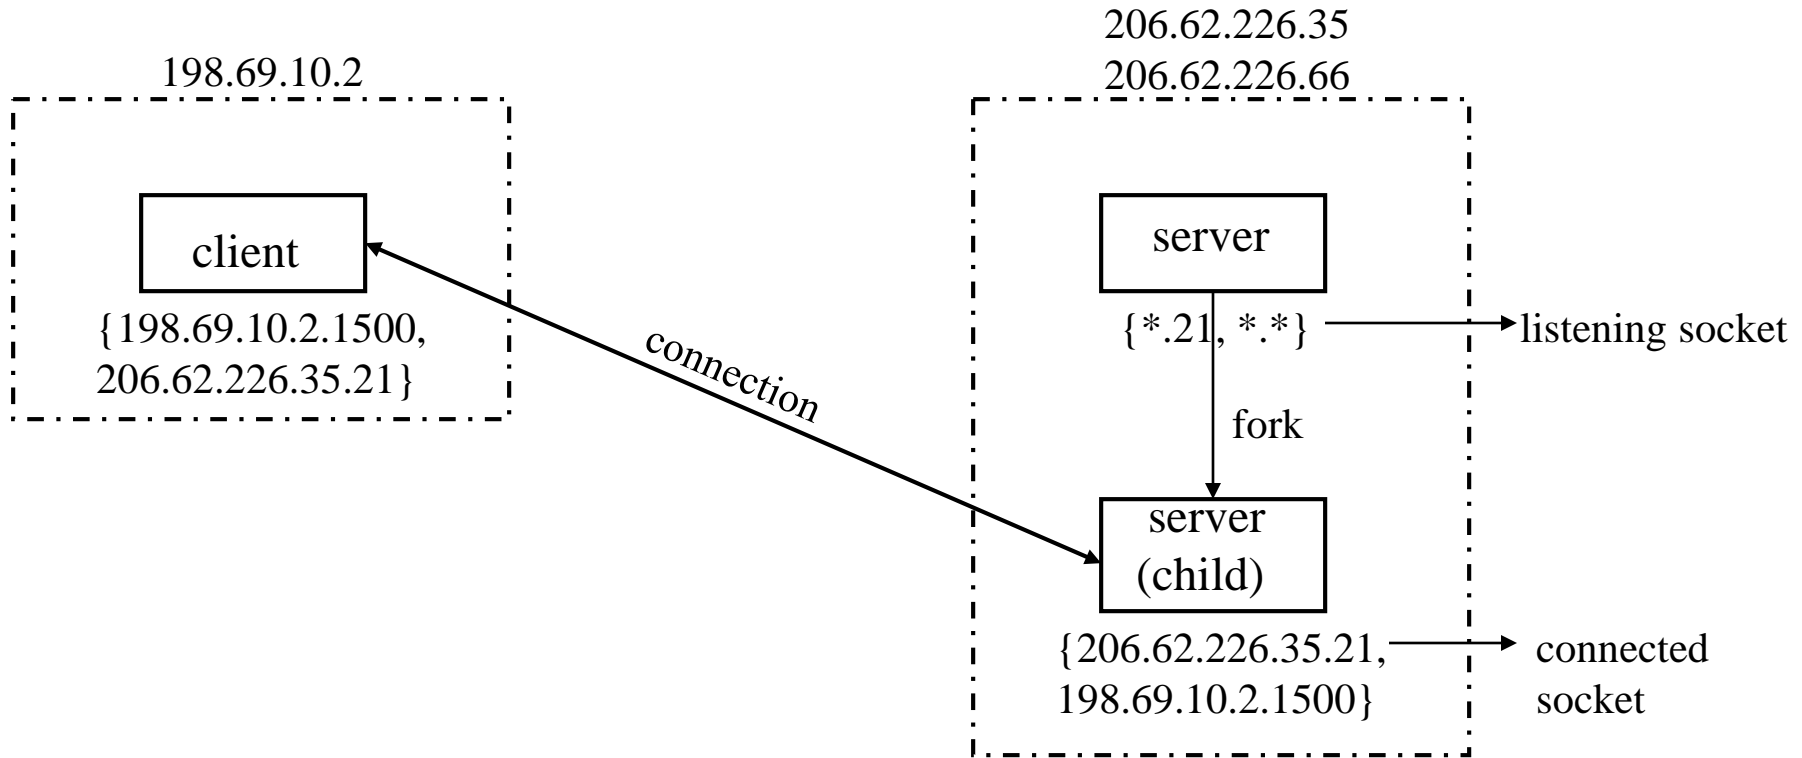


Figure 3. Concurrent server has child handle client.

Connection-oriented Concurrent Server

- `int listenfd, connfd;`
- `if ((listenfd = socket(. . .)) < 0)`
- `err_sys("socket error");`
- `if(bind(listenfd, . . .) < 0)`
- `err_sys("bind error");`
- `if(listen(listenfd, 5) < 0)`
- `err_sys("listen error");`
- `for (; ;) {`
- `connfd = accept(listenfd, . . .); /*blocks */`
- `if (connfd < 0)`
- `err_sys("accept error");`
- `if (fork() == 0) {`
- `close(listenfd);` `/* child */`
- `doit(connfd);`
- `close(connfd);`
- `exit(0);`
- `}`
- `close(connfd);` `/* parent */`
- `}`




Figure 4. Status of client-server before call to *accept*.




Figure 5. Status of client-server after return from *accept*.




Figure 6. Status of client-server after *fork* returns.




Figure 7. Status of client-server after parent and child *close* appropriate sockets.